

11/4/2018

# Machine learning with Energy datasets

## Assignment 2- Report

### Team

ANKIT YADAV  
PREETAM JAIN  
RITU AGARWAL

## **Content**

1. Research
2. Exploratory Data Analysis
3. Feature Engineering
4. Prediction Algorithms
5. Feature Selection
6. Model Validation and Selection
7. Final Pipeline
8. Summary

## 1. Research Paper Analysis

**RESEARCH PAPER 1:**

**RESEARCH PAPER 2:**

**RESEARCH PAPER 3:** Prediction of appliances energy use in smart home

Link - <https://www.sciencedirect.com/science/article/pii/S0360544212002903>

This paper has been written with an aim to predict the energy consumption in household for the next day. To achieve this, they have collected data of homes of France and have analyzed and have come up with certain predictors. It has been studied that residential sector is the biggest sector in electricity consumption. And it is required that we understand the pattern of electricity consumption in household so that Industries can generate and transfer only that amount of energy to the household area to better circulate power. The energy market is divided into distinct categories, but the Day Ahead Market or Spot Market is of great interest. This type of energy market involves bidding the energy consumption of the next day. It is a very complex mechanism, which requires a very good knowledge of the demand for the power suppliers. There were lots of theories which were proposed but it is important to understand the each and every criterion like number of appliances, usage of these appliances, day of week, etc. So, this paper concentrates more over discrimination of usage of electricity on appliance level which would make things easier to understand the pattern of usage of electricity over the course of time. In order to get a better load control, the energy prediction has to go down from total household energy consumption to electrical device consumption. The concept of smart grid has been introduced to tackle power system challenges. Smart grid initiatives seek to improve operations, maintenance and planning using modern technology to better manage energy use and costs. This would help industries to smartly circulate the generated electricity to different industry which would be much more efficient than present method. There have been lot of expectations which was not getting met as the usage of appliance differ over period of time on daily basis. A reliable model was required as usage of appliances on peak time was different as compared to other times. Thus, they came up with a concept called demand dispatch which is ability to control individual loads in precise manner at all the times and not only during peak times. This load management is of two types.

- **Direct Control:** This method refers to classical method of load control which involves increasing the energy production in case of higher load demand.

- **Control by cost:** This method refers to change the load curve shape in such a way that energy consumption peak decreases, even though the total energy consumption for the specific house stays the same. When it has been understood that we have to consider the usage of appliances to get a better picture of electricity consumption and load balancing, there are 4 different type of predictors which can be considered to calculate the same.
- **The “will always consume” predictor:** According to this predictor, we assume that an appliance is always running and consuming electricity.
- **The “will never consume” predictor:** According to this predictor, we assume that an appliance is not at all being used and is not consuming electricity.
- **The ARMA predictor:** ARMA stands for Autoregressive Moving Average. According to this method current value of a time variable is assumed to be a function of its past values and it is expressed as a weighted sum (moving average).
- **The proposed predictor:** According to this model, an inhabitant in the house interacts with various electrical devices as part of his routine activities. Thus, energy consumption can be modeled as a process which is having a random probability distribution or pattern that may be analyzed statistically but may not be predicted precisely.

Improving the precision of prediction is highly necessary. It is important for us to understand the pattern of usage of electricity. The segmentation of data can be made considering various aspects such as the season, month, period of the day (day/night), type of day (weekday/weekend). The objective of this operation is to reduce the average dispersion to improve the prediction. In such conditions, k-means clustering method can be precise to cluster similar data together. At last, I would like to conclude by saying forecasting the energy consumption in homes is an important aspect in the power management of the grid, as the consumption in the residential sector represents a significant percentage in the total electricity demand. The development of the smart grid is not possible without a good prediction of energy consumption. The trend nowadays is to get the prediction of energy consumption not only at house level, but at household appliance level. The prediction of energy consumption in housing is very dependent on inhabitants' behavior, so a stochastic method for prediction has been presented in this paper.

## 2. Exploratory Data Analysis

We imported following libraries

```
In [1]: import numpy as np

In [2]: %matplotlib inline
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
```

Then we read the data and viewed its head.

There were mainly variables for temperature and humidity at various parts of house.

```
In [3]: df=pd.read_csv("../Part_2/energydata_complete.csv")
df.head(10)
df.tail(20)
```

Out[3]:

	date	Appliances	lights	T1	RH_1	T2	RH_2	T3	RH_3	T4	...	T9	RH_9	T_out	Press_n
19715	2016-05-27 14:50:00	60	0	25.500000	46.060000	26.350000	41.000000	28.426667	40.590000	24.666667	...	23.100000	46.590000	21.833333	
19716	2016-05-27 15:00:00	60	0	25.500000	45.933333	26.277143	41.000000	28.356667	40.560000	24.666667	...	23.100000	46.590000	21.800000	
19717	2016-05-27 15:10:00	70	0	25.500000	45.760000	26.200000	41.000000	28.290000	40.433333	24.700000	...	23.133333	46.590000	21.966667	
19718	2016-05-27 15:20:00	80	0	25.500000	45.626667	26.171429	41.000000	28.260000	40.260000	24.700000	...	23.133333	46.590000	22.133333	
19719	2016-05-27 15:30:00	80	0	25.500000	45.590000	26.100000	41.000000	28.200000	40.126667	24.700000	...	23.166667	46.590000	22.300000	
19720	2016-05-27 15:40:00	70	0	25.500000	45.522500	26.100000	41.051429	28.200000	40.200000	24.700000	...	23.100000	46.590000	22.466667	
19721	2016-05-27 15:50:00	100	0	25.500000	45.633333	26.080000	41.196000	28.133333	40.260000	24.700000	...	23.200000	46.590000	22.633333	

We also checked its shape along with the variable information.

```
In [4]: df.shape
```

```
Out[4]: (19735, 29)
```

```
In [8]: data_description = pd.read_csv("../Part_2/variables_description.txt", delimiter='\\t')
data_description.head(40)
```

```
Out[8]:
```

Variable Description	
0	=====
1	date time year-month-day hour:minute:second
2	Appliances, energy use in Wh
3	lights, energy use of light fixtures in the ho...
4	T1, Temperature in kitchen area, in Celsius
5	RH_1, Humidity in kitchen area, in %
6	T2, Temperature in living room area, in Celsius
7	RH_2, Humidity in living room area, in %
8	T3, Temperature in laundry room area
9	RH_3, Humidity in laundry room area, in %
10	T4, Temperature in office room, in Celsius
11	RH_4, Humidity in office room, in %
12	T5, Temperature in bathroom, in Celsius
13	RH_5, Humidity in bathroom, in %
14	T6, Temperature outside the building (outdoor)

We also needed to check the variable information for further use.

```
In [9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19735 entries, 0 to 19734
Data columns (total 29 columns):
date                19735 non-null object
Appliances          19735 non-null int64
lights              19735 non-null int64
T1                  19735 non-null float64
RH_1                19735 non-null float64
T2                  19735 non-null float64
RH_2                19735 non-null float64
T3                  19735 non-null float64
RH_3                19735 non-null float64
T4                  19735 non-null float64
RH_4                19735 non-null float64
T5                  19735 non-null float64
RH_5                19735 non-null float64
T6                  19735 non-null float64
RH_6                19735 non-null float64
T7                  19735 non-null float64
RH_7                19735 non-null float64
T8                  19735 non-null float64
RH_8                19735 non-null float64
T9                  19735 non-null float64
RH_9                19735 non-null float64
T_out               19735 non-null float64
Press_mm_hg         19735 non-null float64
RH_out              19735 non-null float64
```

Then we change the date time according to month, week and time of day.

```

In [11]:
import datetime
import calendar
df['date']=pd.to_datetime(df['date'])
df['year']=df['date'].dt.year
df['month']=df['date'].dt.month
df['day']=df['date'].dt.day
df['day_of_week']=df['date'].dt.weekday_name
df['time_hr_24']=df['date'].dt.hour
morning=range(6,12)
afternoon=range(12,17)
evening=range(17,22)
def time_slot(x):
    if x in morning:
        return 'morning'
    elif x in afternoon:
        return 'afternoon'
    elif x in evening:
        return 'evening'
    else:
        return 'night'

df['day_slot']=df['time_hr_24'].map(time_slot)

week1=range(1,8)
week2=range(8,15)
week3=range(15,22)
week4=range(22,29)

def week_num(x):
    if x in week1:
        return 'week1'
    elif x in week2:
        return 'week2'
    elif x in week3:
        return 'week3'
    elif x in week4:
        return 'week4'
    else:
        return 'week5'

df['week']=df['day'].map(week_num)

```

According to that we added new columns into our data frame

```

df.drop(['date'],axis=1,inplace=True)
df=pd.get_dummies(df,prefix=['DOW','TS','WDT','W'],columns=['day_of_week','day_slot','week_day_type','week'])
print(df.shape)

df.dtypes

```

```

(19735, 50)

```

```

Out[11]: Appliances      int64
lights      int64
T1          float64
RH_1        float64
T2          float64
RH_2        float64
T3          float64
RH_3        float64
T4          float64
RH_4        float64
T5          float64
RH_5        float64
T6          float64
RH_6        float64
T7          float64
RH_7        float64
T8          float64
RH_8        float64
T9          float64
RH_9        float64
T_out       float64
Press_mm_hg float64
RH_out      float64
Windspeed   float64

```

After this we checked the correlation check for redundant variable .

```
In [12]: def get_redundant_pairs(df):
          '''Get diagonal and lower triangular pairs of correlation matrix'''
          pairs_to_drop = set()
          cols = df.columns
          for i in range(0, df.shape[1]):
              for j in range(0, i+1):
                  pairs_to_drop.add((cols[i], cols[j]))
          return pairs_to_drop

          def get_top_abs_correlations(df, n=5):
              au_corr = df.corr().abs().unstack()
              labels_to_drop = get_redundant_pairs(df.select_dtypes(include=['float64', 'int64']))
              au_corr = au_corr.drop(labels=labels_to_drop).sort_values(ascending=False)
              return au_corr[0:n]

          print("Top Absolute Correlations")
          print(get_top_abs_correlations(df.select_dtypes(include=['float64', 'int64']), 20))
```

```
Top Absolute Correlations
rv1    rv2    1.000000
T6     T_out  0.974787
T7     T9     0.944776
T5     T9     0.911055
T3     T9     0.901324
RH_3   RH_4   0.898978
RH_4   RH_7   0.894301
T1     T3     0.892402
T9     month  0.890605
T4     T9     0.889439
T3     T5     0.888169
T1     T5     0.885247
RH_7   RH_8   0.883984
T7     T8     0.882123
RH_1   RH_4   0.880359
T4     T7     0.877763
T1     T4     0.877001
T4     T5     0.871813
T5     T7     0.870624
T8     T9     0.869338
dtype: float64
```

We found that rv1 and rv2, t6 and t\_out etc are highly correlated .

So, we removed these variables from our data.



From the above correlation matrix it could be found:


1. rv1 and rv2 are highly correlated.
2. T-out and T6 are highly correlated.
3. T9 and T7 are highly correlated

Thus, one of them must be rejected.

```
In [19]: df=df.drop(['rv2','T7','T_out'],axis=1)
```

```
In [20]: from matplotlib import cm as cm
```

```
In [21]: # correlation_matrix(df.select_dtypes(include=['float64','int64']))
df.dtypes.value_counts().sort_values().plot(kind='bar')
plt.title('Number of features having certain data type')
plt.show()
```




```
In [22]: fig = plt.figure(figsize = (30,30))
```

We also checked by using heatmap of variables.


```
In [22]: fig = plt.figure(figsize = (30,30))
ax1 = fig.add_subplot(111)
sns.heatmap(df.select_dtypes(include=['float64','int64']).corr())
```

```
Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x21a00ae7438>
```



We also made following plot.


```
In [23]: # df['Total_Usage'] = df['Appliances'] + df['lights']
sns.jointplot(x='T6',y='Appliances',data=df,size=10)
```



We also plotted the graph to look at the changes in temperature over time for various appliances

```
plt.xlabel("Months" , fontsize=15)
plt.ylabel("Mean Temperature in ($^\circ$C)" , fontsize=15)
plt.title("Mean Temperature of every room for first 5 months" , fontsize=20)
y = ['Jan' , 'Feb' , 'Mar' , 'Apr' , 'May']

plt.legend()
plt.show()
```



We also used pandas profiling to check for various variables.

```
In [40]: import pandas_profiling
pandas_profiling.ProfileReport(df)
```

Out[40]:

## Overview

### Dataset info

Number of variables	29
Number of observations	19735
Total Missing (%)	0.0%
Total size in memory	4.4 MiB
Average record size in memory	232.0 B

### Warnings

### Variables types

Numeric	25
Categorical	2
Boolean	0
Date	1
Text (Unique)	0
Rejected	1
Unsupported	0

## Variables

Appliances  
Numeric

Distinct count	92
Unique (%)	0.5%
Missing (%)	0.0%
Missing (n)	0
Infinite (%)	0.0%
Infinite (n)	0

Mean	97.695
Minimum	10
Maximum	1080
Zeros (%)	0.0%



Press\_mm\_hg  
Numeric

Distinct count	2189
Unique (%)	11.1%
Missing (%)	0.0%

Mean	755.52
Minimum	729.3
Maximum	772.3



RH\_2  
Numeric

Distinct count	3376
Unique (%)	17.1%
Missing (%)	0.0%
Missing (n)	0
Infinite (%)	0.0%
Infinite (n)	0

Mean	40.42
Minimum	20.463
Maximum	56.027
Zeros (%)	0.0%




RH\_3  
Numeric

Distinct count	2618
Unique (%)	13.3%
Missing (%)	0.0%
Missing (n)	0

Mean	39.243
Minimum	28.767
Maximum	50.163
Zeros (%)	0.0%



```
In [40]: import pandas_profiling
pandas_profiling.ProfileReport(df)
```



```
In [41]: df.columns
```

### 3. Feature Engineering:

There are lots of features in our dataset. Temperature of 8 different rooms are recorded in degree

Celsius. Humidity of 8 different rooms are recorded in percentage. Outside temperature is also recorded in degree Celsius and humidity in percentage. Along with it pressure has been

recorded in millimeter scale in mercury, visibility in kilometer, dew point in degree Celsius and windspeed in m/s.

Before performing any kind of test, we have to analyse the data and look into it.

We also have to understand how much data does the file contain.

It is important for us to analyze whether there are null values in the dataset and how are values distributed.

```
In [38]: data=pd.read_csv("../Part_2/energydata_complete.csv")
```

```
In [39]: data.head()
```

Out[39]:

	date	Appliances	lights	T1	RH_1	T2	RH_2	T3	RH_3	T4	...	T9	RH_9	T_out	Press_mm_hg	RH_out	Windspeed
0	2016-01-11 17:00:00	60	30	19.89	47.596667	19.2	44.790000	19.79	44.730000	19.000000	...	17.033333	45.53	6.600000	733.5	92.0	7.00
1	2016-01-11 17:10:00	60	30	19.89	46.693333	19.2	44.722500	19.79	44.790000	19.000000	...	17.066667	45.56	6.483333	733.6	92.0	6.66
2	2016-01-11 17:20:00	50	30	19.89	46.300000	19.2	44.626667	19.79	44.933333	18.926667	...	17.000000	45.50	6.366667	733.7	92.0	6.33
3	2016-01-11 17:30:00	50	40	19.89	46.066667	19.2	44.590000	19.79	45.000000	18.890000	...	17.000000	45.40	6.250000	733.8	92.0	6.00
4	2016-01-11 17:40:00	60	40	19.89	46.333333	19.2	44.530000	19.79	45.000000	18.890000	...	17.000000	45.40	6.133333	733.9	92.0	5.66

5 rows x 29 columns



```
In [40]: analysed_data=pd.read_csv("../Part_2/Data_analysed.csv")
```

```
In [41]: analysed_data.head()
```

Out[41]:

Unnamed: 0	date	Appliances	lights	T1	RH_1	T2	RH_2	T3	RH_3	...	month_mean_RH1	month_mean_RH2	month_mean_RH3	...
------------	------	------------	--------	----	------	----	------	----	------	-----	----------------	----------------	----------------	-----

We also did linear regression on the data.

## Linear Regression Model

In [48]:

```
lm=linear_model.LinearRegression()
mod=lm.fit(x_train_sc,y_train)
print(mod.coef_)

print(x_train.columns)

[ 1.27311327e+01 -3.55718214e+00  4.82855496e+01 -3.86638652e+01
 -4.99836910e+01  5.07160517e+01  2.27692226e+01  3.82525914e+00
 -6.79021158e+00 -2.61521677e+00  3.43333957e-01  2.88166878e+01
 -7.34441736e-01  2.44386866e+00 -6.96897198e+00  1.82222358e+01
 -1.37056743e+01 -2.69899999e+01 -4.90771902e+00 -4.09994697e+01
 1.17023936e+00 -6.85337123e+00  3.74400084e+00  2.32604519e+00
 1.85177640e+01  4.09382986e-01  4.09382986e-01 -3.90798505e-14
 -1.17913554e+01 -1.42626392e+00 -1.17279793e+00  1.77942571e+00
 5.20350987e+00  1.13050837e+00  3.87232266e+00 -2.38662247e+00
 -2.52393895e+00 -3.63926792e+00 -1.56566165e+00  1.38663924e+00
 1.01841853e+01  5.50861544e+00 -1.50336310e+01 -1.12071555e+00
 1.12071555e+00]
Index(['lights', 'T1', 'RH_1', 'T2', 'RH_2', 'T3', 'RH_3', 'T4', 'RH_4', 'T5',
       'RH_5', 'T6', 'RH_6', 'T7', 'RH_7', 'T8', 'RH_8', 'T9', 'RH_9', 'T_out',
       'Press_mm_hg', 'RH_out', 'Windspeed', 'Visibility', 'Tdewpoint', 'rv1',
       'rv2', 'year', 'month', 'day', 'time_hr_24', 'time_min', 'DOW_Friday',
       'DOW_Monday', 'DOW_Saturday', 'DOW_Sunday', 'DOW_Thursday',
       'DOW_Tuesday', 'DOW_Wednesday', 'TS_afternoon', 'TS_evening',
       'TS_morning', 'TS_night', 'WDT_weekdays', 'WDT_weekends'],
      dtype='object')
```

Finally, We found the importance of various features in our dataset.

In [77]: `d = pd.DataFrame(feature_importances)`

In [79]: `d.columns = ['Feature', 'Value']`


In [80]: `d`

Out[80]:

	Feature	Value
0	TS_night	0.09
1	T3	0.05
2	time_hr_24	0.05
3	lights	0.04
4	RH_3	0.04
5	RH_5	0.04
6	T8	0.04
7	Press_mm_hg	0.04
8	RH_1	0.03
9	T2	0.03
10	RH_2	0.03
11	T4	0.03
12	RH_4	0.03
13	T7	0.03
14	RH_7	0.03
15	RH_8	0.03
16	RH_9	0.03
17	RH_out	0.03
18	Windspeed	0.03
19	Tdewpoint	0.03
20	T1	0.02

We also plotted a graph to check the same.

```
In [89]: a=d['Feature']
plt.xlabel("Value" , fontsize=20)
plt.ylabel("Feature" , fontsize=20)
plt.legend()
plt.title("Importance of Features" , fontsize=40)
plt.hlines(y=a, xmin=0, xmax=d['Value'], color='skyblue')
plt.plot(d['Value'], a, "o")
plt.rcParams['figure.figsize'] = (25 , 50)
plt.show()
```



In [ ]:

## 4. Prediction Algorithms

```
import pandas as pd
import datetime
import numpy as np
import sklearn
from sklearn.cross_validation import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import RandomizedSearchCV
from sklearn.grid_search import GridSearchCV
from sklearn import linear_model
from sklearn.metrics import *
from sklearn.neural_network import MLPRegressor
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
def mean_absolute_percentage_error(y_true, y_pred):
    y_true, y_pred = np.array(y_true), np.array(y_pred)
    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100
```

### Splitting data

```
In [219]: from sklearn.model_selection import train_test_split
```

```
In [220]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

### Linear Regression

```
In [56]: from sklearn.linear_model import LinearRegression
```

```
In [57]: lm=LinearRegression()
```

```
In [218]: X = df.drop(['Appliances'],axis=1)
y = energy_data_complete['Appliances']
```



```
In [221]: lm.fit(X_train,y_train)

Out[221]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
              normalize=False)

In [225]: # The coefficients
print('Coefficients: \n', lm.coef_)

Coefficients:
[ 1.58639142e+00 -1.81851521e+00  1.32188205e+01 -1.78769245e+01
 -1.27339698e+01  2.57171413e+01  6.65504364e+00  1.44641205e+00
 -1.46278964e+00 -1.82584390e+00  1.37015448e-02  3.31368293e+00
 -1.26498195e-01 -1.24450234e+00  1.02147795e+01 -2.88283445e+00
 -1.35422575e+01 -1.16098376e+00  1.89036791e-01  9.85002954e-01
 1.88603123e+00  1.75320013e-01 -2.13444209e+00  4.24421146e-02
 -3.55271368e-15 -8.81883689e+00  6.49793896e-01 -1.92191360e-01
 1.46465722e+01  3.74873514e+00  1.05233084e+01 -7.92293360e+00
 -6.89346415e+00 -9.13086101e+00 -4.97135705e+00  1.51364888e+00
 2.19771521e+01  1.11796838e+01 -3.46704849e+01 -2.60037485e+00
 2.60037485e+00  8.58817418e+00  6.60952869e+00  2.53724874e+00
 -7.25238248e+00 -1.04825691e+01]
```

## Using Random Forest

```
In [231]: from sklearn.ensemble import RandomForestRegressor

In [232]: # Random Forest Regressor
# model = RandomForestRegressor(n_estimators=300, max_features = 11)
# # create the RFE model and select 3 attributes
# rfe = RFE(model)
# rfe = rfe.fit(X_train, y_train)
# # summarize the selection of the attributes
# print(rfe.support_)
# print(rfe.ranking_)
# print(rfe.n_features_)
# #Check the accuracy of the model
# rfe.score(X_train, y_train)
rand_forest_model = RandomForestRegressor(max_depth=5, random_state=0) # build model
rand_forest_model = rand_forest_model.fit(X_train, y_train.values.ravel()) # train model
RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=5,
                      max_features='auto', max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
                      oob_score=False, random_state=0, verbose=0, warm_start=False)
feature_importance = rand_forest_model.feature_importances_
R2 = rand_forest_model.score(X_train, y_train) # coefficient of determination
rand_forest_model.predictions = rand_forest_model.predict(X_test) # make predictions

C:\Users\ritua\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\ensemble\forest.py:248: FutureWarning: The default v
alue of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
  "10 in version 0.20 to 100 in 0.22.", FutureWarning)

In [233]: # rand_forest_model_metrics = view_metrics(y_test, rand_forest_model_predictions, 'Random Forest Regressor', samples = 50, total_s
plt.scatter(y_test,rand_forest_model_predictions)
plt.xlabel('Y Test')
plt.ylabel('Predicted Y')


Out[233]: Text(0,0.5,'Predicted Y')
```

```
In [234]: print('MAE:', metrics.mean_absolute_error(y_test, rand_forest_model_predictions))
print('MSE:', metrics.mean_squared_error(y_test, rand_forest_model_predictions))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, rand_forest_model_predictions)))
```

```
MAE: 48.96025500004364
MSE: 7862.571568228872
RMSE: 88.67114281562448
```

```
In [235]: sns.distplot((y_test-rand_forest_model_predictions),bins=50)
```

```
Out[235]: <matplotlib.axes._subplots.AxesSubplot at 0x1fa2c04d320>
```



## Neural Network Models

### Neural Network Model

```
In [22]: mlp = MLPRegressor(hidden_layer_sizes=(365,365,365),max_iter=500,alpha=1.0000000e-06,random_state=42)
mlp.fit(x_train_sc,y_train)
```

```
Out[22]: MLPRegressor(activation='relu', alpha=1e-06, batch_size='auto', beta_1=0.9,
beta_2=0.999, early_stopping=False, epsilon=1e-08,
hidden_layer_sizes=(365, 365, 365), learning_rate='constant',
learning_rate_init=0.001, max_iter=500, momentum=0.9,
nesterovs_momentum=True, power_t=0.5, random_state=42, shuffle=True,
solver='adam', tol=0.0001, validation_fraction=0.1, verbose=False,
warm_start=False)
```

### Neural Network on Training Dataset

```
In [23]: y_train_pred=mlp.predict(x_train_sc)
print("R2 : ",r2_score(y_train,y_train_pred))
print("MAE : ",mean_absolute_error(y_train,y_train_pred))
print("RMSE : ",np.sqrt(mean_squared_error(y_train,y_train_pred)))
print("MAPE : ",mean_absolute_percentage_error(y_train,y_train_pred))
```

```
R2 : 0.7592006178700873
MAE : 27.93967498755577
RMSE : 50.52214265006127
MAPE : 30.235642921597133
```

### Neural Network on Testing Dataset

```
In [24]: y_test_pred=mlp.predict(x_test_sc)
print("R2 : ",r2_score(y_test,y_test_pred))
print("MAE : ",mean_absolute_error(y_test,y_test_pred))
print("RMSE : ",np.sqrt(mean_squared_error(y_test,y_test_pred)))
print("MAPE : ",mean_absolute_percentage_error(y_test,y_test_pred))
```

```
R2 : 0.4078561865573932
MAE : 39.471749093105366
RMSE : 78.10655339647808
MAPE : 39.30815270527815
```

## **5. Feature Selection**

In machine learning and statistics, feature selection, also known as variable selection, attribute selection or variable subset selection, is the process of selecting a subset of relevant features (variables, predictors) for use in model construction. Feature selection techniques are used for four reasons:

- simplification of models to make them easier to interpret by researchers/users,
- shorter training times,
- enhanced generalization by reducing overfitting (formally, reduction of variance)

The central premise when using a feature selection technique is that the data contains many features that are either redundant or irrelevant, and can thus be removed without incurring much loss of information. Redundant or irrelevant features are two distinct notions, since one relevant feature may be redundant in the presence of another relevant feature with which it is strongly correlated.

### **TPOT:**

The Tree-Based Pipeline Optimization Tool (TPOT) was one of the very first AutoML methods and open-source software packages developed for the data science community. The goal of TPO

is to automate the building of ML pipelines by combining a flexible expression tree representation of pipelines with stochastic search algorithms such as genetic programming. TPOT makes use of the Python-based scikit-learn library as its ML menu

## Tpot

```
In [251]: from tpot import TPOTRegressor  
pipeline_optimizer = TPOTRegressor()
```

```
In [252]: pipeline_optimizer = TPOTRegressor(generations=10, population_size=20, cv = 3,  
                                             random_state=42, verbosity=2)
```

```
In [253]: pipeline_optimizer.fit(X_train,y_train)
```

Failed to display Jupyter Widget of type `HBox`.

If you're reading this message in the Jupyter Notebook or JupyterLab Notebook, it may mean that the widgets JavaScript is still loading. If this message persists, it likely means that the widgets JavaScript library is either not installed or not enabled. See the [Jupyter Widgets Documentation](#) for setup instructions.

If you're reading this message in another frontend (for example, a static rendering on GitHub or [NBViewer](#)), it may mean that your frontend doesn't currently support widgets.

```
Generation 1 - Current best internal CV score: -5225.879046662546  
Generation 2 - Current best internal CV score: -5225.879046662546  
Generation 3 - Current best internal CV score: -5222.948696227385  
Generation 4 - Current best internal CV score: -5222.948696227385  
Generation 5 - Current best internal CV score: -5222.948696227385  
Generation 6 - Current best internal CV score: -5222.948696227385  
Generation 7 - Current best internal CV score: -5222.948696227385  
Generation 8 - Current best internal CV score: -5222.948696227385  
Generation 9 - Current best internal CV score: -5222.948696227385  
Generation 10 - Current best internal CV score: -5222.948696227385
```

```
Best pipeline: ElasticNetCV(RandomForestRegressor(input_matrix, bootstrap=False, max_features=0.25, min_samples_leaf=1, min_samples_split=2, n_estimators=100), l1_ratio=0.05, tol=0.01)
```

```
Out[253]: TPOTRegressor(config_dict=None, crossover_rate=0.1, cv=3,  
                        disable_update_check=False, early_stop=None, generations=10,  
                        max_eval_time_mins=5, max_time_mins=None, memory=None,  
                        mutation_rate=0.9, n_jobs=1, offspring_size=None,  
                        periodic_checkpoint_folder=None, population_size=20,  
                        random_state=42, scoring=None, subsample=1.0, use_dask=False,  
                        verbosity=2, warm_start=False)
```

## Info of model

```

In [254]: energy_data_complete.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19735 entries, 0 to 19734
Data columns (total 27 columns):
date                19735 non-null datetime64[ns]
Appliances          19735 non-null int64
lights              19735 non-null int64
T1                  19735 non-null float64
RH_1                19735 non-null float64
T2                  19735 non-null float64
RH_2                19735 non-null float64
T3                  19735 non-null float64
RH_3                19735 non-null float64
T4                  19735 non-null float64
RH_4                19735 non-null float64
T5                  19735 non-null float64
RH_5                19735 non-null float64
T6                  19735 non-null float64
RH_6                19735 non-null float64
RH_7                19735 non-null float64
T8                  19735 non-null float64
RH_8                19735 non-null float64
T9                  19735 non-null float64
RH_9                19735 non-null float64
Press_mm_hg         19735 non-null float64
RH_out              19735 non-null float64
Windspeed           19735 non-null float64
Visibility           19735 non-null float64
Tdewpoint           19735 non-null float64
rv1                 19735 non-null float64
month               19735 non-null int64
dtypes: datetime64[ns](1), float64(23), int64(3)
memory usage: 4.1 MB

```

When we get the output, we check the accuracy of score.

```

In [255]: print(pipeline_optimizer.score(X_test, y_test))

-3693.3057863344193

```

```

In [ ]:

```

## TSFresh

tsfresh is a python package. It automatically calculates a large number of time series characteristics, the so called features. Further the package contains methods to evaluate the explaining power and importance of such characteristics for regression or classification tasks.

We start by creating the model for testing.

Performing tsfresh

```
In [9]: from tsfresh.utilities.dataframe_functions import roll_time_series
```

```
In [10]: df_shift, y = make_forecasting_frame(x, kind="price", max_timeshift=10, rolling_direction=1)
df_shift
```

Out[10]:

	time	value	id	kind
154008	2016-01-11 21:30:00	100.0	2016-01-11 21:40:00	price
136892	2016-01-11 21:30:00	100.0	2016-01-11 21:50:00	price
154009	2016-01-11 21:40:00	100.0	2016-01-11 21:50:00	price
119777	2016-01-11 21:30:00	100.0	2016-01-11 22:00:00	price
136893	2016-01-11 21:40:00	100.0	2016-01-11 22:00:00	price
154010	2016-01-11 21:50:00	100.0	2016-01-11 22:00:00	price
102663	2016-01-11 21:30:00	100.0	2016-01-11 22:10:00	price
119778	2016-01-11 21:40:00	100.0	2016-01-11 22:10:00	price
136894	2016-01-11 21:50:00	100.0	2016-01-11 22:10:00	price
154011	2016-01-11 22:00:00	110.0	2016-01-11 22:10:00	price
85550	2016-01-11 21:30:00	100.0	2016-01-11 22:20:00	price
102664	2016-01-11 21:40:00	100.0	2016-01-11 22:20:00	price
119779	2016-01-11 21:50:00	100.0	2016-01-11 22:20:00	price
136895	2016-01-11 22:00:00	110.0	2016-01-11 22:20:00	price
154012	2016-01-11 22:10:00	400.0	2016-01-11 22:20:00	price
68438	2016-01-11 21:30:00	100.0	2016-01-11 22:40:00	price
85551	2016-01-11 21:40:00	100.0	2016-01-11 22:40:00	price

```
In [12]: y.head()
```

```
Out[12]: date
2016-01-11 21:40:00    100
2016-01-11 21:50:00    100
2016-01-11 22:00:00    110
2016-01-11 22:10:00    400
2016-01-11 22:20:00    400
Name: value, dtype: int64
```

```
In [13]: from tsfresh import extract_relevant_features
```

Extracting relevant features.

```
In [14]: from tsfresh import select_features
from tsfresh.utilities.dataframe_functions import impute

impute(X)
features_filtered = select_features(X, y)

WARNING:tsfresh.feature_selection.relevance:Inferred classification as machine learning task
```

```
In [15]: features_filtered.head()
```

2016-01-11 21:40:00	91.421879	100.0	111.541685	100.0	87.
2016-01-11 21:50:00	91.421879	100.0	111.541685	100.0	87.
2016-01-11 22:00:00	91.421879	100.0	111.541685	100.0	87.
2016-01-11 22:10:00	91.421879	100.0	111.541685	100.0	87.
2016-01-11 22:20:00	91.421879	100.0	111.541685	100.0	87.

Getting final accuracy.

```
# write obtained data
TestModels = TestModels.append([tmp])

TestModels.set_index('Model', inplace=True)

print(TestModels)
```

	MAE_Test	MAE_Train	MSE_Test	MSE_Train \
Model				
LinearRegression	2.622187e+01	2.409460e+01	6.085721e+03	2.077187e+03
RandomForestRegressor	2.465078e+01	9.751281e+00	2.382641e+03	4.056849e+02
MLPRegressor	1.188270e+08	3.977830e+08	3.407596e+19	6.007033e+20

	R2_Test	R2_Train	RMSE_Test	RMSE_Train
Model				
LinearRegression	-4.474209e-01	5.351238e-01	7.801103e+01	4.557617e+01
RandomForestRegressor	4.333154e-01	9.092074e-01	4.881230e+01	2.014162e+01
MLPRegressor	-8.104589e+15	-1.344379e+17	5.837462e+09	2.450925e+10

In [ ]:

## Boruta

An all relevant feature selection wrapper algorithm. It finds relevant features by comparing original attributes' importance with importance achievable at random, estimated using their permuted copies (shadows).

First, build the model.

Building Model

```
In [2]: df_train, df_test = train_test_split(df, train_size=0.7, random_state=42)
x_train = df_train.iloc[:, 1:]
y_train = df_train['Appliances']
scaler = scaler.fit(x_train)
x_train_sc = scaler.transform(x_train)
x_test = df_test.iloc[:, 1:]
y_test = df_test['Appliances']
x_test_sc = scaler.transform(x_test)
```

Running Boruta on this model

Running the boruta.

```
In [4]: import pandas as pd
# from sklearn.ensemble import RandomForestClassifier
from boruta import BorutaPy

# Load X and y
# NOTE BorutaPy accepts numpy arrays only, hence the .values attribute
X = x_train_sc
y = y_train

# define random forest classifier, with utilising all cores and
# sampling in proportion to y labels
rf = RandomForestRegressor(n_jobs=-1, max_depth=25)

# define Boruta feature selection method
feat_selector = BorutaPy(rf, n_estimators='auto', verbose=2)

# find all relevant features
feat_selector.fit(X, y)
```

```
Iteration:    1 / 100
Confirmed:    0
Tentative:    45
Rejected:     0
Iteration:    2 / 100
Confirmed:    0
Tentative:    45
Rejected:     0
Iteration:    3 / 100
Confirmed:    0
Tentative:    45
Rejected:     0
Iteration:    4 / 100
Confirmed:    0
Tentative:    45
Rejected:     0
Iteration:    5 / 100
Confirmed:    0
Tentative:    45
```

We need to select only helping features from our model.



We have to select only few features from our dataset which would help us give better results in our prediction.

```
In [10]: # check selected features
feat_selector.support_

Out[10]: array([ True, False,  True,  True,  True,  True,  True,  True,  True,
        True,  True,  True, False, False,  True,  True,  True,  True,
        True, False,  True,  True, False, False, False, False, False,
        False, False,  True, False, False, False, False, False,
        False, False, False, False, False, False,  True, False, False])
```

Implementing the optimised features which we got from above step after performing boruta.

```
In [11]: column_list = ['lights', 'RH_1', 'T2', 'RH_2', 'T3', 'RH_3', 'T4', 'RH_4', 'T5', 'RH_5', 'T6', 'RH_7', 'T8', 'RH_8', 'T9', 'RH_9', 'Press_mm_hg', '']
x_train=df_train.iloc[:,1:]
x_train= x_train[column_list]
print(x_train.shape)
y_train=df_train['Appliances']
scaler.fit(x_train)
x_train_sc=scaler.transform(x_train)
x_test=df_test.iloc[:,1:]
x_test = x_test[column_list]
print(x_test.shape)
y_test=df_test['Appliances']
x_test_sc=scaler.transform(x_test)

(11982, 20)
(5136, 20)
```

## Making a model for random forest

```
In [12]: rf=RandomForestRegressor(bootstrap=False, criterion='mse', max_depth=25,
    max_features='sqrt', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=350, n_jobs=1,
    oob_score=False, random_state=42, verbose=0, warm_start=False)
rf.fit(x_train_sc,y_train)

Out[12]: RandomForestRegressor(bootstrap=False, criterion='mse', max_depth=25,
    max_features='sqrt', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=350, n_jobs=1,
    oob_score=False, random_state=42, verbose=0, warm_start=False)
```

## Checking the accuracy of training and testing data

Checking the accuracy of data for training set.

```
In [13]: y_train_pred=rf.predict(x_train_sc)
print("R2 : ",r2_score(y_train,y_train_pred))
print("MAE : ",mean_absolute_error(y_train,y_train_pred))
print("RMSE : ",np.sqrt(mean_squared_error(y_train,y_train_pred)))
print("MAPE : ",mean_absolute_percentage_error(y_train,y_train_pred))

R2 : 0.999365728603491
MAE : 0.978722739191654
RMSE : 1.6866280335790647
MAPE : 1.5898191622381659
```

Checking the accuracy of data for testing set.

```
In [14]: y_test_pred=rf.predict(x_test_sc)
print("R2 : ",r2_score(y_test,y_test_pred))
print("MAE : ",mean_absolute_error(y_test,y_test_pred))
print("RMSE : ",np.sqrt(mean_squared_error(y_test,y_test_pred)))
print("MAPE : ",mean_absolute_percentage_error(y_test,y_test_pred))

R2 : 0.5493916369664105
MAE : 21.217272013210984
RMSE : 43.3163158169478
MAPE : 25.665825612557814
```

## 6. Model Validation and Selection

Creating a model

```
In [1]: import pandas as pd
import datetime
import numpy as np
import sklearn
from sklearn.cross_validation import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import RandomizedSearchCV
from sklearn.grid_search import GridSearchCV
from sklearn import linear_model
from sklearn.metrics import *
from sklearn.neural_network import MLPRegressor
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
import matplotlib.pyplot as plt
%matplotlib inline
from matplotlib.pylab import rcParams
rcParams['figure.figsize'] = 6, 4
def mean_absolute_percentage_error(y_true, y_pred):
    y_true, y_pred = np.array(y_true), np.array(y_pred)
    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100
```


Now the model is ready, let's find out how the data is distributed.

```
In [2]: energy.head()
```

Out[2]:

	Appliances	lights	T1	RH_1	T2	RH_2	T3	RH_3	T4	RH_4	...	DOW_Sunday	DOW_Thursday	DOW_Tuesday
27	100	20	21.356667	45.826667	20.666667	45.163333	20.390000	46.090000	19.390000	47.500000	...	0	0	
28	100	20	21.390000	45.690000	20.700000	45.060000	20.390000	46.090000	19.426667	47.993333	...	0	0	
29	100	20	21.500000	45.333333	20.700000	44.933333	20.390000	46.060000	19.566667	48.466667	...	0	0	
30	110	20	21.500000	45.126667	20.790000	44.633333	20.390000	46.000000	19.666667	48.093333	...	0	0	
31	400	20	21.533333	44.966667	20.790000	44.360000	20.426667	45.933333	19.600000	47.500000	...	0	0	

5 rows x 46 columns



Now when we now how the initial data is distributed, let's find out whether it has any null values. Also, we have to find what are the nature of the data in our dataset.

```
In [3]: energy.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 17118 entries, 27 to 19697
Data columns (total 46 columns):
Appliances      17118 non-null int64
lights          17118 non-null int64
T1              17118 non-null float64
RH_1            17118 non-null float64
T2              17118 non-null float64
RH_2            17118 non-null float64
T3              17118 non-null float64
RH_3            17118 non-null float64
T4              17118 non-null float64
RH_4            17118 non-null float64
DOW_Sunday      17118 non-null int64
DOW_Thursday    17118 non-null int64
DOW_Tuesday     17118 non-null int64
DOW_Wednesday   17118 non-null int64
DOW_Friday      17118 non-null int64
DOW_Saturday    17118 non-null int64
```

Detecting overfitting.

```
In [4]: # Overfitting detecting
from sklearn.preprocessing import PolynomialFeatures
from sklearn import linear_model
from sklearn.model_selection import train_test_split
from sklearn.metrics import *
```

Making the model to help us better analyze the data better.

```
In [5]: df_train, df_test = train_test_split(df, train_size=0.7, random_state=42)
x_train = df_train.iloc[:, 1:]
y_train = df_train['Appliances']
scaler = scaler.fit(x_train)
X_train = scaler.transform(x_train)
x_test = df_test.iloc[:, 1:]
y_test = df_test['Appliances']
X_test = scaler.transform(x_test)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_split.py:2026: FutureWarning: From version 0.21, test_size
will always complement train_size unless both are specified.
FutureWarning)
```

## Regularization


```
In [6]: # Set random_state as a const to make sure the same split will be generated every time you run the code.
#X_train, X_test, y_train, y_test = train_test_split(X, Y, random_state=0)

# Containers for note down the MSE
train_mse_list = []
test_mse_list = []
degree_of_polynomial = []
```

```
In [7]: # MSE of quartic expression will be very large compared with the lower degree, so let's just end at 3.
```

```
for i in range(0, 4):
    model = PolynomialFeatures(degree=i)
    # Transfer the X to a polynomial form by using fit_transform
    X_train_ = model.fit_transform(X_train)
    X_test_ = model.fit_transform(X_test)
    lm = linear_model.LinearRegression()
    lm.fit(X_train_, y_train)
    train_pred = lm.predict(X_train_)
    train_mse_list.append(mean_squared_error(y_train, train_pred))
    test_pred = lm.predict(X_test_)
    test_mse_list.append(mean_squared_error(y_test, test_pred))
    degree_of_polynomial.append(i)
    print("\nDegree : ", i)
    print("For Training Data : ")
    print("R2 : ", r2_score(y_train, train_pred))
    print("MAE : ", mean_absolute_error(y_train, train_pred))
    print("RMSE : ", np.sqrt(mean_squared_error(y_train, train_pred)))
    print("MAPE : ", mean_absolute_percentage_error(y_train, train_pred))
    print("\nFor Testing Data : ")
    print("R2 : ", r2_score(y_test, test_pred))
    print("MAE : ", mean_absolute_error(y_test, test_pred))
    print("RMSE : ", np.sqrt(mean_squared_error(y_test, test_pred)))
    print("MAPE : ", mean_absolute_percentage_error(y_test, test_pred))
```

```
plt.xlabel('degree of polynomial')
plt.ylabel('MSE')
plt.grid(True)
plt.title('Polynomial Regression without Regularization')
plt.plot(degree_of_polynomial, train_mse_list, '-b', degree_of_polynomial, test_mse_list, '-g')
plt.show()
```



## Performing Lasso

Performing L1 (lasso)

```
In [8]: from sklearn.linear_model import Lasso


# Randomly pick a alpha value for regularization
l1reg = Lasso(alpha=0.000003, normalize=True)

l1reg_test_mse_list = []

for i in range(0, 4):
    model = PolynomialFeatures(degree=i)
    X_train_ = model.fit_transform(X_train)
    X_test_ = model.fit_transform(X_test)

    l1reg.fit(X_train_, y_train)
    train_pred_l1 = l1reg.predict(X_train_)
    test_pred_l1 = l1reg.predict(X_test_)
    l1reg_test_mse_list.append(mean_squared_error(y_test, test_pred_l1))
    print("\nDegree : ", i)
    print("For Training Data : ")
    print("R2 : ", r2_score(y_train, train_pred_l1))
    print("MAE : ", mean_absolute_error(y_train, train_pred_l1))
    print("RMSE : ", np.sqrt(mean_squared_error(y_train, train_pred_l1)))
    print("MAPE : ", mean_absolute_percentage_error(y_train, train_pred_l1))
    print("\nFor Testing Data : ")
    print("R2 : ", r2_score(y_test, test_pred_l1))
    print("MAE : ", mean_absolute_error(y_test, test_pred_l1))
    print("RMSE : ", np.sqrt(mean_squared_error(y_test, test_pred_l1)))
    print("MAPE : ", mean_absolute_percentage_error(y_test, test_pred_l1))

plt.xlabel('degree of polynomial')
plt.ylabel('MSE')
plt.grid(True)
plt.title('Polynomial Regression without L1(Lasso) Regularization')
plt.plot(degree_of_polynomial, test_mse_list, '-b', degree_of_polynomial, l1reg_test_mse_list, '--r')
plt.show()
```



## Cross Validation Technique

Testing for training dataset.

```
In [15]: #df_train,df_test = train_test_split(df,train_size=0.7,random_state=42)
#x_train=df_train.iloc[:,1:]
#y_train=df_train['Appliances']
#scaler.fit(x_train)
#X_train=scaler.transform(x_train)
#x_test=df_test.iloc[:,1:]
#y_test=df_test['Appliances']
#X_test=scaler.transform(x_test)
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LinearRegression
accuracy_train = cross_val_score(estimator = LinearRegression() , X = X_train, y = y_train , cv = 10)
accuracy_train
```

```
Out[15]: array([0.23939319, 0.20994031, 0.28080389, 0.24029031, 0.24235254,
0.21612244, 0.23405993, 0.24911729, 0.24016138, 0.26527946])
```

Mean accuracy.

```
In [16]: accuracy_train.mean()
```

```
Out[16]: 0.24175207193743198
```

Performing on testing dataset.

```
In [17]: from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LinearRegression
accuracy_test = cross_val_score(estimator = LinearRegression() , X = X_test, y = y_test , cv = 10)
accuracy_test
```

```
Out[17]: array([0.15211076, 0.20928208, 0.25724492, 0.20743369, 0.26784327,
0.25292395, 0.28916533, 0.22152727, 0.18031048, 0.22997972])
```

## Getting Accuracy of model

```
In [18]: accuracy_test.mean()
```

```
Out[18]: 0.2267821480096499
```

```
In [ ]:
```

## 7 Final Pipeline

```
In [238]: from sklearn.preprocessing import MinMaxScaler  
min_max=MinMaxScaler()
```

```
In [239]: X_train_minmax = min_max.fit_transform(X_train)
```

```
C:\Users\ritua\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\preprocessing\data.py:323: DataConversionWarning: Data with input dtype uint8, int64, float64 were all converted to float64 by MinMaxScaler.  
    return self.partial_fit(X, y)
```

```
In [240]: X_test_minmax=min_max.fit_transform(X_test)
```

```
C:\Users\ritua\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\preprocessing\data.py:323: DataConversionWarning: Data with input dtype uint8, int64, float64 were all converted to float64 by MinMaxScaler.  
    return self.partial_fit(X, y)
```

```
In [241]: from sklearn.preprocessing import StandardScaler
```

```
In [242]: scaler = StandardScaler()
```

```
In [243]: X_train_scale=scaler.fit_transform(X_train)
```

```
C:\Users\ritua\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\preprocessing\data.py:617: DataConversionWarning: Data with input dtype uint8, int64, float64 were all converted to float64 by StandardScaler.  
    return self.partial_fit(X, y)  
C:\Users\ritua\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\base.py:462: DataConversionWarning: Data with input dtype uint8, int64, float64 were all converted to float64 by StandardScaler.  
    return self.fit(X, **fit_params).transform(X)
```

Error metric

```

In [245]: error_metric = pd.DataFrame({'r2_train': [],
                                      'r2_test': [],
                                      'rms_train': [],
                                      'rms_test': [],
                                      'mae_train': [],
                                      'mae_test': [],
                                      'mape_train': [],
                                      'mape_test': []})

rmse_dict = {}

def calc_error_metric(modelname, model, X_train_scale, y_train, X_test_scale, y_test):
    global error_metric
    y_train_predicted = model.predict(X_train)
    y_test_predicted = model.predict(X_test)

    #MAE, RMS, MAPE, R2

    r2_train = r2_score(y_train, y_train_predicted)
    r2_test = r2_score(y_test, y_test_predicted)

    rms_train = sqrt(mean_squared_error(y_train, y_train_predicted))
    rms_test = sqrt(mean_squared_error(y_test, y_test_predicted))

    mae_train = mean_absolute_error(y_train, y_train_predicted)
    mae_test = mean_absolute_error(y_test, y_test_predicted)

    mape_train = np.mean(np.abs((y_train - y_train_predicted) / y_train)) * 100
    mape_test = np.mean(np.abs((y_test - y_test_predicted) / y_test)) * 100

    rmse_dict[modelname] = rms_test

    df_local = pd.DataFrame({'Model': [modelname],
                              'r2_train': [r2_train],
                              'r2_test': [r2_test],
                              'rms_train': [rms_train],
                              'rms_test': [rms_test],
                              'mae_train': [mae_train],
                              'mae_test': [mae_test],
                              'mape_train': [mape_train],
                              'mape_test': [mape_test]})

```

Getting best model.



```
In [246]: from sklearn.neural_network import MLPRegressor
```

```
In [247]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
In [248]: from math import sqrt
```

```
In [249]: import operator
```

```
In [250]: # Regression
clf = LinearRegression()
clf.fit(X_train, y_train)
calc_error_metric('Linear Regression', clf, X_train, y_train, X_test, y_test)
print('Linear Regression Peformed')

# Random Forest
rf = RandomForestRegressor(n_estimators=100, max_depth=7)
rf.fit(X_train, y_train)
calc_error_metric('RandomForest', rf, X_train, y_train, X_test, y_test)
print('RandomForest Performed')

# Neural network
nn = MLPRegressor()
nn.fit(X_train, y_train)
calc_error_metric('Neural Network', nn, X_train, y_train, X_test, y_test)
print('Neural Network performed')

#### Calculate best model
best_model = min(rmse_dict.items(),key=operator.itemgetter(1))[0]
print('Best Model is-', best_model)

### Write the error
error_metric.to_csv('Error_metrics.csv')

Linear Regression Peformed
RandomForest Performed
Neural Network performed
Best Model is  RandomForest
```

```
In [90]: error_metric
```


```
Out[90]:
```

	Model	mae_test	mae_train	mape_test	mape_train	r2_test	r2_train	rms_test	rms_train
0	Linear Regression	52.491344	52.625376	62.606520	60.304493	0.178557	0.189465	90.333257	93.139558
0	RandomForest	45.879759	44.400356	51.149829	47.843062	0.272879	0.359836	84.988958	82.773972
0	Neural Network	51.305092	51.746769	57.898395	55.929704	0.124915	0.122877	93.236107	96.889885

Getting feature importance of random forest

```
In [91]: # Feature Importance in Random Forest
imp_feature = pd.Series(rf.feature_importances_,X_train.columns).sort_values(ascending=False)
imp_feature.plot(kind='bar', title='Feature Importances',figsize = (15,8))
plt.ylabel('Feature Importance Score')
```

Out[91]: Text(0,0.5,'Feature Importance Score')



Doing prediction by random forest.


```
In [231]: from sklearn.ensemble import RandomForestRegressor
```

```
In [232]: # Random Forest Regressor
# model = RandomForestRegressor(n_estimators=300, max_features = 11)
# # create the RFE model and select 3 attributes
# rfe = RFE(model)
# rfe = rfe.fit(X_train, y_train)
# # summarize the selection of the attributes
# print(rfe.support_)
# print(rfe.ranking_)
# print(rfe.n_features_)
# #Check the accuracy of the model
# rfe.score(X_train, y_train)
rand_forest_model = RandomForestRegressor(max_depth=5, random_state=0) # build model
rand_forest_model = rand_forest_model.fit(X_train, y_train.values.ravel()) # train model
RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=5,
    max_features='auto', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
    oob_score=False, random_state=0, verbose=0, warm_start=False)
feature_importance = rand_forest_model.feature_importances_
R2 = rand_forest_model.score(X_train, y_train) # coefficient of determination
rand_forest_model_predictions = rand_forest_model.predict(X_test) # make predictions
```

C:\Users\ritua\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\ensemble\forest.py:248: FutureWarning: The default value of n\_estimators will change from 10 in version 0.20 to 100 in 0.22.  
"10 in version 0.20 to 100 in 0.22.", FutureWarning)

```
In [233]: # rand_forest_model_metrics = view_metrics(y_test, rand_forest_model_predictions, 'Random Forest Regressor', samples = 50, total_s
plt.scatter(y_test, rand_forest_model_predictions)
plt.xlabel('Y Test')
plt.ylabel('Predicted Y')
```


Out[233]: Text(0,0.5,'Predicted Y')



```
In [234]: print('MAE:', metrics.mean_absolute_error(y_test, rand_forest_model_predictions))
print('MSE:', metrics.mean_squared_error(y_test, rand_forest_model_predictions))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, rand_forest_model_predictions)))
```

MAE: 48.96025500004364  
MSE: 7862.571568228872  
RMSE: 88.67114281562448

```
Out[235]: <matplotlib.axes._subplots.AxesSubplot at 0x1fa2c04d320>
```



### Predicting using Random forest algorithms on train data

```
In [236]: from sklearn.feature_selection import RFE
```

```
In [237]: model = RandomForestRegressor(n_estimators=20, max_features = 11)
rfe = RFE(model)
rfe = rfe.fit(X_train, y_train)
# summarize the selection of the attributes
print(rfe.support_)
print(rfe.ranking_)
print(rfe.n_features_)
#Check the accuracy of the model
rfe.score(X_train, y_train)
```

```
[ True True True True True True True True True True True True True
 True True True True True True True True True True True True True
 False False False True False False False False False False False
 False False True False False False False False False False False]
[ 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
 24 7 5 1 8 9 11 15 23 17 21 12 6 13 1 19 10 18 22 14 20 16]
```

Out[237]: 0.9224397516383712

Now, using the same on our testing data.

```
In [90]: model = RandomForestRegressor(n_estimators=20, max_features = 11)
rfe = RFE(model)
rfe = rfe.fit(X_train, y_train)
# summarize the selection of the attributes
print(rfe.support_)
print(rfe.ranking_)
print(rfe.n_features_)
#Check the accuracy of the model
rfe.score(X_test, y_test)
```

```
[ True False  True  True  True  True  True  True  True  True  True  True
  True  True  True  True  True  True  True  True  True  True  True  True
  False False  False  True  False False  False False False False False False
  False False  True  False False  False  False False False False False]
[ 1  2  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  3  6  1  1
 24  7  4  1  8  9 10 10 21 11 22 13  5 17  1 12 15 20 16 18 14 23]
```

Out[90]: 0.5297603171841764

**So, our model gives 52.9 % accuracy.**

## **8. Summary**

By the above analysis we have concluded following points regarding the data given to us.

- Best model to analyze and predict is Random Forest.
- There are many columns which are highly correlated, and they need to be removed in order to get good prediction
- The data have almost none outliers and no NULL valued column. So, the data is almost clean.